

# MovieLens Report

Akram Zalloum  
December 2021

## Introduction

This recommendation system is used to predict the ratings that will be assigned to the movies based on the historical users' ratings. Based on the competition announced by Netflix in 2006 to solve a challenge in building a model used to enhance the prediction on users' ratings and predict what rate each movie will get. This happened and the model enhance the prediction by 10% which worth a lot of money, in 2009 winner announced who build this model. We will try to reflect this in below report.

## Data set

The MovieLens Data set collected by GroupLens Research available in (<http://movielens.org>). The data will be split in training and testing sets where 10% will be used for validation and evaluation.

```
1 # Create edx set, validation set (final hold-out test set)
2 #####
3 # Note: this process could take a couple of minutes
4 if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
5 if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
6 if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")
7 library(tidyverse)
8 library(caret)
9 library(data.table)
10 # MovieLens 10M dataset:
11 # https://grouplens.org/datasets/movielens/10m/
12 # http://files.grouplens.org/datasets/movielens/ml-10m.zip
13 dl <- tempfile()
14 download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)
15 ratings <- fread(text = gsub(":", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
16                 col.names = c("userId", "movieId", "rating", "timestamp"))
17 movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\:::", 3)
18 colnames(movies) <- c("movieId", "title", "genres")
19 # if using R 3.6 or earlier:
20 movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(levels(movieId))[movieId],
21                                         title = as.character(title),
22                                         genres = as.character(genres))
23 # if using R 4.0 or later:
24 movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(movieId),
25                                         title = as.character(title),
26                                         genres = as.character(genres))
27 movielens <- left_join(ratings, movies, by = "movieId")
28 # Validation set will be 10% of MovieLens data
29 set.seed(1, sample.kind="Rounding") # if using R 3.5 or earlier, use `set.seed(1)`
30 test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
31 edx <- movielens[-test_index,]
32 temp <- movielens[test_index,]
33 # Make sure userId and movieId in validation set are also in edx set
34 validation <- temp %>%
35   semi_join(edx, by = "movieId") %>%
36   semi_join(edx, by = "userId")
37 # Add rows removed from validation set back into edx set
38 removed <- anti_join(temp, validation)
39 edx <- rbind(edx, removed)
40 rm(dl, ratings, movies, test_index, temp, movielens, removed)
```

## Understand data

Based on the data structure we have 6 columns (userId,movieID,rating,timestamp,title,genres).

```
str(edx)
```

```
## 'data.frame': 9000055 obs. of 6 variables:
## $ userId : int 1 1 1 1 1 1 1 1 1 1 ...
## $ movieId : num 122 185 292 316 329 355 356 362 364 370 ...
## $ rating : num 5 5 5 5 5 5 5 5 5 5 ...
## $ timestamp: int 838985046 838983525 838983421 838983392 838983392 838983392 838984474 838983653 838984885
## $ title : chr "Boomerang (1992)" "Net, The (1995)" "Outbreak (1995)" "Stargate (1994)" ...
## $ genres : chr "Comedy|Romance" "Action|Crime|Thriller" "Action|Drama|Sci-Fi|Thriller" "Action
```

```
summary(edx)
```

```
##      userId      movieId      rating      timestamp
## Min.   : 1      Min.   : 1      Min.   :0.500      Min.   :7.897e+08
## 1st Qu.:18124    1st Qu.: 648    1st Qu.:3.000    1st Qu.:9.468e+08
## Median :35738    Median : 1834    Median :4.000    Median :1.035e+09
## Mean   :35870    Mean   : 4122    Mean   :3.512    Mean   :1.033e+09
## 3rd Qu.:53607    3rd Qu.: 3626    3rd Qu.:4.000    3rd Qu.:1.127e+09
## Max.   :71567    Max.   :65133    Max.   :5.000    Max.   :1.231e+09
##      title      genres
## Length:9000055    Length:9000055
## Class :character    Class :character
## Mode :character    Mode :character
##
##
```


Regarding the ratings statistics we have Min = 1, Max = 5, Mean = 3.51 and Mode = 4.0.

```
## # A tibble: 5 x 2
##   rating count
##   <dbl> <int>
## 1     4 2588430
## 2     3 2121240
## 3     5 1390114
## 4    3.5 791624
## 5     2 711422
```


To get number of unique movies and users, we use the code:

```
##   n_users n_movies
## 1   69878   10677
```


The histogram below to get number of ratings by movie



To get number of ratings by user



Visualization below to show the rating by movie genre



Will list below the top 10 genres

```
## # A tibble: 20 x 2
##   genres      count
##   <chr>      <int>
## 1 Drama    3910127
## 2 Comedy   3540930
## 3 Action    2560545
## 4 Thriller  2325899
## 5 Adventure 1908892
## 6 Romance   1712100
## 7 Sci-Fi    1341183
## 8 Crime     1327715
## 9 Fantasy   925637
## 10 Children 737994
## 11 Horror    691485
## 12 Mystery   568332
## 13 War       511147
## 14 Animation 467168
## 15 Musical   433080
## 16 Western   189394
## 17 Film-Noir 118541
## 18 Documentary 93066
## 19 IMAX       8181
## 20 (no genres listed) 7
```

Prepare training and testing data sets (80% to 20%)

```
1 set.seed(1)
2 test_index <- createDataPartition(y = edx$rating, times = 1, p = 0.2, list = FALSE)
3 train_set <- edx[-test_index,]
4 test_set <- edx[test_index,]
```

Calculate RMSE to measure the accuracy of the model.

```
1 RMSE <- function(true_ratings, predicted_ratings){
2   sqrt(mean((true_ratings - predicted_ratings)^2, na.rm = TRUE))
3 }
```

## Model 1

Predict the same rating for all movies regardless of the user and without any constraints.

```
1 mu_1 <- mean(train_set$rating)
2 mu_1
```

```
## [1] 3.512482
```

```
1 naive_rmse <- RMSE(test_set$rating,mu_1)
2 naive_rmse
```

```
## [1] 1.059909
```


Get the RMSE and compare results.

```
1 rmse_results <- data_frame(method = "The average", RMSE = naive_rmse)
2 rmse_results %>% knitr::kable()
```

Method	RMSE
The average	1.0599094

## Model 2

Where some movies rated more than others and based on the previous model will add  $b_i$  to represent the average ranking for movie  $i$  to show the biased, so we compute the average.




```
1 predicted_ratings <- Mu_2 + test_set %>%
2   left_join(movie_avgs, by='movieId') %>%
3   pull(b_i)
4 model_2_rmse <- RMSE(predicted_ratings, test_set$rating)
5 rmse_results <- bind_rows(rmse_results,
6   data_frame(method="Movie Effect Model",
7     RMSE = model_2_rmse))
8 rmse_results %>% knitr::kable()
```

Method	RMSE
The average	1.0599094
Movie Effect Model	0.9437429

### Model 3

Users who rate too many movies (>100)



Based on the clear gaps between ratings will get the mean

```
1 user_avgs <- train_set %>%
2   left_join(movie_avgs, by='movieId') %>%
3   group_by(userId) %>%
4   summarize(b_u = mean(rating - Mu_2 - b_i))
```

Will calculate RMSE again to see the enhancement

```
1 predicted_ratings <- test_set %>%
2   left_join(movie_avgs, by='movieId') %>%
3   left_join(user_avgs, by='userId') %>%
4   mutate(pred = Mu_2 + b_i + b_u) %>%
5   pull(pred)
6 model_3_rmse <- RMSE(predicted_ratings, test_set$rating)
7 rmse_results <- bind_rows(rmse_results,
8                           data_frame(method="Movie + User Effects Model",
9                                       RMSE = model_3_rmse))
10 rmse_results %>% knitr::kable()
```

Method	RMSE
The average	1.0599094
Movie Effect Model	0.9437429
Movie + User Effects Model	0.8659320

## Apply RMSE for validation

```
1 valid_pred_rating <- validation %>%
2   left_join(movie_avgs , by = "movieId" ) %>%
3   left_join(user_avgs , by = "userId") %>%
4   mutate(pred = Mu_2 + b_i + b_u ) %>%
5   pull(pred)
6 model_3_valid <- RMSE(validation$rating, valid_pred_rating)
7 rmse_results <- bind_rows( rmse_results, data_frame(Method = "Validation Results" , RMSE = model_3_valid))
8 rmse_results%>% knitr::kable()
```

Method	RMSE	Method
The average	1.0599094	NA
Movie Effect Model	0.9437429	NA
Movie + User Effects Model	0.8659320	NA
NA	0.8664515	Validation Results

## Findings

Based on the methodology followed which is naive the 3rd model has best RMSE, after using the linear regression the findings that old movies have better prediction.